

Apache processes and threads in mod_ndb




mod_ndb.h

```
struct mod_ndb_process {  
    unsigned short n_connections;  
    unsigned short n_threads;  
    struct mod_ndb_connection conn; // not a pointer  
};
```





mod_ndb.h

```
struct mod_ndb_connection {  
    unsigned int connected;  
    int ndb_force_send;  
    Ndb_cluster_connection *connection;  
    ndb_instance **instances;  
    struct mod_ndb_connection *next;  
};  
typedef struct mod_ndb_connection ndb_connection;
```

One mod_ndb_process
per Apache process



0 --> n_connections
One mod_ndb_connection per NDB connect string



n_threads

One ndb_instance per Apache thread,
per NDB connect string

mod_ndb.h

```
struct mod_ndb_instance {  
    struct mod_ndb_connection *conn;  
    Ndb *db;  
    NdbTransaction *tx;  
    unsigned int requests;  
    unsigned int errors;  
    unsigned int declined;  
    unsigned int row_not_found;  
    unsigned int row_found;  
};  
typedef struct mod_ndb_instance  
    ndb_instance;
```

Using C++ class templates above the Apache API

Apache's C-language API relies heavily on void pointers that you can cast to different data types. In C++, though, casting is no fun – the compiler requires you to make every cast explicitly, and casting defeats the type-safe design of the language.

Here are some examples from the array API: `array_header->elts` is a `char *` which you cast to an array pointer, and `ap_push_array()` returns a void pointer to a new element.

httpd/ap_alloc.h

```
typedef struct {
    ap_pool *pool;
    int elt_size;
    int nelts;
    int nalloc;
    char *elts;
} array_header;

array_header * ap_make_array(pool *p, int nelts, int elt_size);
void * ap_push_array(array_header *);
```

```
template <class T>
class apache_array: public array_header {
public:
    int size() { return this->nelts; }
    T **handle() { return (T**) &(this->elts); }
    T *items() { return (T*) this->elts; }
    T &item(int n){ return ((T*) this->elts)[n]; }
    T *new_item() { return (T*) ap_push_array(this); }
    void * operator new(size_t, ap_pool *p, int n) {
        return ap_make_array(p, n, sizeof(T));
    };
};
```

mod_ndb.h

In `mod_ndb`, the template `apache_array<T>` builds a subclass of `array_header` to manage an array of any type. All of the casting is done here in the template definition, so the code in the actual source files is cleaner:

```
dir->visible      = new(p, 4) apache_array<char *>;
dir->updatable    = new(p, 4) apache_array<char *>;
dir->indexes      = new(p, 2) apache_array<config::index>;

*dir->visible->new_item() = ap_pstrdup(cmd->pool, arg);
```

Per-server (i.e. per-VHOST) config structure

config::srv
connect_string

```
namespace config {
    /* Apache per-server configuration */
    struct srv {
        char *connect_string; };
}
```

Apache per-directory config structure

config::dir
database
table
pathinfo_size
pathinfo
allow_delete
results
format_param[]
visible
updatable
indexes
key_columns

```
/* Apache per-directory configuration */
namespace config {
    struct dir {
        char *database;
        char *table;
        int pathinfo_size;
        short *pathinfo;
        int allow_delete;
        result_format results;
        char *format_param[2];
        apache_array<char*> *visible;
        apache_array<char*> *updatable;
        apache_array<config::index> *indexes;
        apache_array<config::key_col> *key_columns;
    };
}
```

Configuration Directives

*mod_ndb.cc and
config.cc*

Directive	Function	Data Structure	Inheritable
ndb-connectstring	ap_set_string_slot()	srv->connect_string	Yes
Database	ap_set_string_slot()	dir->database	Yes
Table	ap_set_string_slot()	dir->table	Yes
Deletes	ap_set_flag_slot()	dir->allow_delete	Yes
Format	result_format()	dir->results	Yes
Columns	non_key_column()	dir->visible	No
AllowUpdate	non_key_column()	dir->updatable	No
PrimaryKey	primary_key()	dir->key_columns	No
UniqueIndex	named_index()	dir->key_columns	No
OrderedIndex	named_index()	dir->key_columns	No
PathInfo	pathinfo()	dir->pathinfo	No
Filter	filter()	dir->key_columns	No

encoding and decoding
of NDB & MySQL data types

MySQL
Time()
Date()
Datetime()
String()
result()
value()

```
namespace MySQL {
    char *Time(pool *p, const NdbRecAttr &rec);
    char *Date(pool *p, const NdbRecAttr &rec);
    char *Datetime(pool *p, const NdbRecAttr &rec);
    char *String(pool *p, const NdbRecAttr &rec, enum ndb_string_packing packing);
    char *result(pool *p, const NdbRecAttr &rec);
    mvalue value(pool *p, const NdbDictionary::Column *col, const char *val);
};
```

- Time(), Date() and Datetime() decode specially packed mysql data types.

- String() can unpack three different sorts of strings packed into NDB character arrays.

```
enum ndb_string_packing {
    char_fixed,
    char_var,
    char_longvar
};
```

- result() is a generic "decode" function; it converts an NdbRecAttr to a printable ASCII value

- value() is a generic "encode" function; given an ASCII value (from HTTP) and an NdbDictionary::Column (which specifies how to encode the value, it will return an mvalue properly encoded for the database

```
struct mvalue {
    mvalue_use use_value;
    union {
        const char *    val_const_char;
        char *          val_char;
        int              val_signed;
        unsigned int     val_unsigned;
        long long        val_64;
        unsigned long long val_unsigned_64;
        float            val_float;
        double           val_double;
        const NdbDictionary::Column * err_col;
    } u;
};
typedef struct mvalue mvalue;
```

mod_ndb Architecture: Formatting of Results

Results can be formatted in a variety of ways

```
mod_ndb.h
enum result_format
{
    json = 1,
    raw,
    xml,
    ap_note
}
```

JSON Result Formatting

```
class JSON {
public:
    static const char * new_array;      // "[\n"
    static const char * end_array;      // "]\n"
    static const char * new_object;     // "{\n"
    static const char * end_object;     // "}\n"
    static const char * delimiter;      // ", "
    static const char * is              // ": "
    static char *value(const NdbRecAttr &rec, request_rec *r);
    inline static char *member(const NdbRecAttr &rec, request_rec *r) {
        return ap_pstrcat(r->pool,
                           rec.getColumn()->getName(),
                           JSON::is,
                           JSON::value(rec,r),
                           NULL);
    }
};
```

JSON.h

MySQL
Time()
Date()
Datetime()
String()
result()
value()

JSON
new_array
end_array
new_object
end_object
delimiter
is
value()
member()

JSON::value() is largely a wrapper around MySQL::result(), but strings, dates, and times are all quoted, and NULLs are represented as `"null"`

Indexes and key columns

config::index
name
type
n_columns
first_col_serial
first_col_idx

config::key_col
name
is_in_pk
is_filter
filter_op
index_id
serial_no
idx_map_bucket
filter_col_serial
filter_col
next_in_key_serial
next_in_key

/*

Every time a new column is added, the columns get reshuffled some, so we have to fix all the mappings between serial numbers and actual column id numbers.

The configuration API in Apache never gives the module a chance to "finalize" a configuration structure. You never know when you're finished with a particular directory. So, we run `fix_all_columns()` every time we create a new column, which, alas, does not scale too well.

While processing the config file, the CPU time spent fixing columns grows with n^2 , the square of the number of columns. This could be improved using config handling that was more complex (a container directive) or less user-friendly (an explicit "end" token).

On the other hand, the design is optimized for handling queries at runtime, where some operations (e.g. following the list of columns that belong to an index) are constant, and the worst (looking up a column name in the columns table) grows at $\log n$.

*/